

```
1 !pip install kaggle

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: kaggle in /usr/local/lib/python3.9/dist-packages (1.5.13)
Requirement already satisfied: requests in /usr/local/lib/python3.9/dist-packages (from kaggle) (2.27.1)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.9/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.9/dist-packages (from kaggle) (8.0.1)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.9/dist-packages (from kaggle) (1.26.15)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.9/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.9/dist-packages (from kaggle) (4.65.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.9/dist-packages (from kaggle) (2022.12.7)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.9/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.9/dist-packages (from requests->kaggle) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.9/dist-packages (from requests->kaggle) (3.4)

1 from google.colab import files
2 files.upload()

Choose Files kaggle.json
• kaggle.json(application/json) - 65 bytes, last modified: 3/28/2023 - 100% done
Saving kaggle.json to kaggle.json
{'kaggle.json':
b'{"username": "zakenmani" "key": "7757c32ff01f25a734c3a9c265c3afa" }'

1 !mkdir ~/.kaggle
2 !cp kaggle.json ~/.kaggle/

mkdir: cannot create directory '/root/.kaggle': File exists

1 !chmod 600 ~/.kaggle/kaggle.json

Saved successfully!
abshshahane/fake-news-classification

Downloading fake-news-classification.zip to /content
98% 90.0M/92.1M [00:03<00:00, 27.5MB/s]
100% 92.1M/92.1M [00:03<00:00, 28.0MB/s]

1 !unzip fake-news-classification.zip

Archive: fake-news-classification.zip
  inflating: WELFake_Dataset.csv

1 import pandas as pd
2
3 df = pd.read_csv('WELFake_Dataset.csv')

1 df

   Unnamed: 0  title  text  label
0            0  LAW ENFORCEMENT ON HIGH ALERT Following Threat...  No comment is expected from Barack Obama Membe...  1
1            1      NaN  Did they post their votes for Hillary already?  1
2            2  UNBELIEVABLE! OBAMA'S ATTORNEY GENERAL SAYS MO...  Now, most of the demonstrators gathered last ...  1
3            3  Bobby Jindal, raised Hindu, uses story of Chri...  A dozen politically active pastors came here f...  0
4            4  SATAN 2: Russia unvelis an image of its terrif...  The RS-28 Sarmat missile, dubbed Satan 2, will...  1
...         ...      ...      ...      ...
72129        72129  Russians steal research on Trump in hack of U...  WASHINGTON (Reuters) - Hackers believed to be ...  0
72130        72130  WATCH: Giuliani Demands That  You know, because in  1

1 # get cols of the df
2 df.columns
```

```

1 # sentiment is the y value, review is the x value. We split with an 80/20 split
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import LabelEncoder
4 import numpy as np
5 # Filter out any rows with NaN values
6 df = df.dropna()
7
8 # Separate features and target variable
9 X = df[['title', 'text']]
10 y = df['label']
11
12 # Convert categorical variables using Label Encoding
13 label_encoder = LabelEncoder()
14 X['title'] = label_encoder.fit_transform(X['title'])
15 X['text'] = label_encoder.fit_transform(X['text'])
16 y = label_encoder.fit_transform(y)
17
18 # Split data into training and testing sets
19 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
20
21 # make sure that all text is string
22 X_train = X_train.astype(str)
23 X_test = X_test.astype(str)
24 y_train = np.asarray(y_train).astype('float32').reshape((-1,1))
25 y_test = np.asarray(y_test).astype('float32').reshape((-1,1))

```

<ipython-input-13-531550a5e943>:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

Saved successfully!

ation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

<ipython-input-13-531550a5e943>:15: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
X['text'] = label_encoder.fit_transform(X['text'])

```

1 import seaborn as sns
2
3 sns.set_style('darkgrid')
4 sns.countplot(x='label', data=df).set(title='Label Distribution for Training Data')

```

[Text(0.5, 1.0, 'Label Distribution for Training Data')]



Describe the data set and what the model should be able to predict.

This data set contains real news and fake news, with titles and text being assigned to each. It contains 4 columns, but as the first is just an identifier, we only consider the other 3. These are: Title (about the text news heading); Text (about the news content); and Label (0 = fake and 1 = real).

The model should be able to accurately predict whether a given entry is real or fake news given the title and text of the entry. There is an 80/20 split between training/testing data for the model to predict on.

```

1 # Sequential
2
3 from tensorflow import keras
4 from keras.models import Sequential
5 from keras.layers import Dense, Embedding, Flatten
6 from keras.preprocessing.text import Tokenizer
7 from keras.utils import to_categorical
8 from tensorflow.keras.preprocessing.sequence import pad_sequences
9 import pandas as pd
10
11 # Tokenize and pad sequences
12 max_words = 10000
13 tokenizer = Tokenizer(num_words=max_words)
14 tokenizer.fit_on_texts(X_train['title'] + ' ' + X_train['text'])
15 X_train_seq = tokenizer.texts_to_sequences(X_train['title'] + ' ' + X_train['text'])
16 X_test_seq = tokenizer.texts_to_sequences(X_test['title'] + ' ' + X_test['text'])
17 max_seq_length = max(len(x) for x in X_train_seq)
18 X_train_padded = pad_sequences(X_train_seq, maxlen=max_seq_length)
19 X_test_padded = pad_sequences(X_test_seq, maxlen=max_seq_length)
20
21 # Define model
22 model = Sequential()
23 model.add(Embedding(max_words, 100, input_length=max_seq_length))
24 model.add(Dense(100, activation='relu'))
25 model.add(Dense(100, activation='relu'))
26 model.add(Dense(32, activation='relu'))
27 model.add(Dense(1, activation='sigmoid'))
28
29 # Compile model
30 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
31
32 # Train model
33 model.fit(X_train_padded, y_train, epochs=15, batch_size=16, validation_data=(X_test_padded, y_test))

```

Epoch 1/15
3577/3577 [=====] - 32s 8ms/step - loss: 1.8217 - accuracy: 0.6226 - val_loss: 0.7168 - val_accuracy: 0.5239
Epoch 2/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6434 - accuracy: 0.6485 - val_loss: 0.6881 - val_accuracy: 0.5553
Epoch 3/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6503 - accuracy: 0.6500 - val_loss: 0.6974 - val_accuracy: 0.5553
Epoch 4/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6418 - accuracy: 0.6500 - val_loss: 0.6761 - val_accuracy: 0.5553
Epoch 5/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6416 - accuracy: 0.6500 - val_loss: 0.6939 - val_accuracy: 0.5553
Epoch 6/15
3577/3577 [=====] - 18s 5ms/step - loss: 0.6415 - accuracy: 0.6498 - val_loss: 0.6778 - val_accuracy: 0.5553
Epoch 7/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6404 - accuracy: 0.6500 - val_loss: 0.6831 - val_accuracy: 0.5553
Epoch 8/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6406 - accuracy: 0.6500 - val_loss: 0.6826 - val_accuracy: 0.5553
Epoch 9/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6402 - accuracy: 0.6500 - val_loss: 0.6782 - val_accuracy: 0.5553
Epoch 10/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6418 - accuracy: 0.6500 - val_loss: 0.6849 - val_accuracy: 0.5553
Epoch 11/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6402 - accuracy: 0.6500 - val_loss: 0.6825 - val_accuracy: 0.5553
Epoch 12/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6402 - accuracy: 0.6500 - val_loss: 0.6843 - val_accuracy: 0.5553
Epoch 13/15
3577/3577 [=====] - 16s 4ms/step - loss: 0.6420 - accuracy: 0.6500 - val_loss: 0.6799 - val_accuracy: 0.5553
Epoch 14/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6401 - accuracy: 0.6500 - val_loss: 0.6862 - val_accuracy: 0.5553
Epoch 15/15
3577/3577 [=====] - 17s 5ms/step - loss: 0.6401 - accuracy: 0.6500 - val_loss: 0.6829 - val_accuracy: 0.5553
<keras.callbacks.History at 0x7f1e7bda5eb0>

```

1 # CNN
2
3 from keras.layers import Dense, Embedding, Conv1D, MaxPooling1D, GlobalMaxPooling1D
4

```

```

5 # Tokenize and pad sequences; higher max word count and seq length
6 max_words = 10000
7 tokenizer = Tokenizer(num_words=max_words)
8 tokenizer.fit_on_texts(X_train['title'] + ' ' + X_train['text'])
9 X_train_seq = tokenizer.texts_to_sequences(X_train['title'] + ' ' + X_train['text'])
10 X_test_seq = tokenizer.texts_to_sequences(X_test['title'] + ' ' + X_test['text'])
11 max_seq_length = 1000 # Increase the max sequence length
12 X_train_padded = pad_sequences(X_train_seq, maxlen=max_seq_length)
13 X_test_padded = pad_sequences(X_test_seq, maxlen=max_seq_length)
14
15 # Define model
16 model = Sequential()
17 model.add(Embedding(max_words, output_dim=50, input_length=max_seq_length, mask_zero=True))
18 model.add(Conv1D(32, 5, activation='relu'))
19 model.add(MaxPooling1D(5))
20 model.add(Conv1D(64, 5, activation='relu'))
21 model.add(MaxPooling1D(5))
22 model.add(Conv1D(128, 5, activation='relu'))
23 model.add(MaxPooling1D(5))
24 model.add(Conv1D(256, 5, activation='relu'))
25 model.add(GlobalMaxPooling1D())
26 model.add(Dense(1, activation='sigmoid'))
27
28 # Compile model
29 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
30
31 # Train model
32 model.fit(X_train_padded, y_train, epochs=15, batch_size=16, validation_data=(X_test_padded, y_test))

Epoch 1/15
3577/3577 [=====] - 32s 8ms/step - loss: 0.6273 - accuracy: 0.6048 - val_loss: 0.6899 - val_accuracy: 0.5082
Saved successfully! [=====] - 30s 8ms/step - loss: 0.5879 - accuracy: 0.6291 - val_loss: 0.6933 - val_accuracy: 0.5083
3577/3577 [=====] - 29s 8ms/step - loss: 0.5881 - accuracy: 0.6292 - val_loss: 0.6896 - val_accuracy: 0.5083
Epoch 4/15
3577/3577 [=====] - 32s 9ms/step - loss: 0.5876 - accuracy: 0.6293 - val_loss: 0.6923 - val_accuracy: 0.5083
Epoch 5/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5878 - accuracy: 0.6293 - val_loss: 0.6895 - val_accuracy: 0.5083
Epoch 6/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5876 - accuracy: 0.6293 - val_loss: 0.6958 - val_accuracy: 0.5083
Epoch 7/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6903 - val_accuracy: 0.5083
Epoch 8/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6929 - val_accuracy: 0.5083
Epoch 9/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6931 - val_accuracy: 0.5083
Epoch 10/15
3577/3577 [=====] - 30s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6908 - val_accuracy: 0.5083
Epoch 11/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6932 - val_accuracy: 0.5083
Epoch 12/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5877 - accuracy: 0.6293 - val_loss: 0.6928 - val_accuracy: 0.5083
Epoch 13/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6936 - val_accuracy: 0.5083
Epoch 14/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6926 - val_accuracy: 0.5083
Epoch 15/15
3577/3577 [=====] - 28s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6917 - val_accuracy: 0.5083
<keras.callbacks.History at 0x7f1dcaeedf0>

1 # Embeddings
2
3 from gensim.models import Word2Vec
4
5 # Concatenate title and text columns
6 X_train_concat = X_train['title'] + ' ' + X_train['text']
7 X_test_concat = X_test['title'] + ' ' + X_test['text']
8
9 # Tokenize sequences
10 max_words = 10000
11 tokenizer = Tokenizer(num_words=max_words)
12 tokenizer.fit_on_texts(X_train_concat)
13 X_train_seq = tokenizer.texts_to_sequences(X_train_concat)
14 X_test_seq = tokenizer.texts_to_sequences(X_test_concat)
15
16 # Pad sequences
17 max_seq_length = 1000
18 X_train_padded = pad_sequences(X_train_seq, maxlen=max_seq_length)

```

```

19 X_test_padded = pad_sequences(X_test_seq, maxlen=max_seq_length)
20
21 # Train Word2Vec model
22 embedding_dim = 50
23 w2v_model = Word2Vec(sentences=X_train_seq, window=5, min_count=1, workers=4)
24
25 # Initialize embedding matrix with zeros
26 embedding_matrix = np.zeros((max_words, embedding_dim))
27
28 # Build embedding matrix
29 for word, i in tokenizer.word_index.items():
30     if i < max_words:
31         if word in w2v_model.wv.key_to_index:
32             embedding_matrix[i] = w2v_model.wv[word]
33
34
35 # Define model
36 model = Sequential()
37 model.add(Embedding(max_words, output_dim=50, input_length=max_seq_length, mask_zero=True))
38 model.add(Conv1D(32, 5, activation='relu'))
39 model.add(MaxPooling1D(5))
40 model.add(Conv1D(64, 5, activation='relu'))
41 model.add(MaxPooling1D(5))
42 model.add(Conv1D(128, 5, activation='relu'))
43 model.add(MaxPooling1D(5))
44 model.add(Conv1D(256, 5, activation='relu'))
45 model.add(GlobalMaxPooling1D())
46 model.add(Dense(1, activation='sigmoid'))
47
48
49 # Compile model
50 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
51
52 # Train model
53 model.fit(X_train_padded, y_train, epochs=15, batch_size=16, validation_data=(X_test_padded, y_test))

```

Epoch 1/15
3577/3577 [=====] - 33s 9ms/step - loss: 0.6377 - accuracy: 0.5902 - val_loss: 0.6875 - val_accuracy: 0.5083
Epoch 2/15
3577/3577 [=====] - 28s 8ms/step - loss: 0.5901 - accuracy: 0.6272 - val_loss: 0.6941 - val_accuracy: 0.5083
Epoch 3/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5881 - accuracy: 0.6292 - val_loss: 0.6921 - val_accuracy: 0.5084
Epoch 4/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5875 - accuracy: 0.6293 - val_loss: 0.6936 - val_accuracy: 0.5083
Epoch 5/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5875 - accuracy: 0.6293 - val_loss: 0.6913 - val_accuracy: 0.5083
Epoch 6/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5875 - accuracy: 0.6293 - val_loss: 0.6924 - val_accuracy: 0.5083
Epoch 7/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6917 - val_accuracy: 0.5083
Epoch 8/15
3577/3577 [=====] - 28s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6956 - val_accuracy: 0.5083
Epoch 9/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5875 - accuracy: 0.6293 - val_loss: 0.6946 - val_accuracy: 0.5083
Epoch 10/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6902 - val_accuracy: 0.5083
Epoch 11/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6921 - val_accuracy: 0.5083
Epoch 12/15
3577/3577 [=====] - 27s 7ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6945 - val_accuracy: 0.5083
Epoch 13/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6909 - val_accuracy: 0.5083
Epoch 14/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6911 - val_accuracy: 0.5083
Epoch 15/15
3577/3577 [=====] - 27s 8ms/step - loss: 0.5874 - accuracy: 0.6293 - val_loss: 0.6927 - val_accuracy: 0.5083
<keras.callbacks.History at 0x7f1df255ba00>

Analysis

I tried three different approaches when it came to classification of real/fake news - Sequential, CNN, and embedding.

For all 3, they used 15 epochs and 16 batch size.

Sequential

Sequential was the simplest model, but had the highest accuracy out of the three. The end accuracy was 0.65, and the highest of the lot. There were no embedding layers, only Dense.

CNN

I opted for a CNN over an RNN as this data isn't sequential/time series, thus an RNN isn't proper use case for this. This yielded an accuracy of 0.6293, and used the MaxPooling functions for the purposes of the network.

Embedding

The Embedded CNN used the Word2vec embedding to learn word associations. After training the w2v model, the weights of this model were used when applying it to the actual CNN. This also yielded an accuracy of 0.6293, tying it with CNN and below Sequential.

Conclusion

Overall, I had a good time working with the different models to try and figure out what was going on with each during this process. However, I was surprised at the Sequential performance beating out the other models.

Saved successfully!



✓ 7m 26s completed at 4:33 PM

